

GPU Introduction

JSC OpenACC Course 2023

24 October 2023 | Andreas Herten | Forschungszentrum Jülich

Outline

Introduction

GPU History

GPU History

Architecture Comparison

Jülich Systems

JUWELS Cluster

JUWELS Booster

JURECA DC

App Showcase

Platform

3 Core Features

Memory

Asynchronicity

SIMT

Generation Comparison

High Throughput

Summary

Programming GPUs

Libraries

GPU programming models

CUDA C/C++

Parallel Model

Conclusions

History of GPUs

A short but unparalleled story

- 1999 Graphics computation pipeline implemented in dedicated *graphics hardware*
 - Computations using OpenGL graphics library [2]
 - »GPU« coined by NVIDIA [3]

History of GPUs

A short but unparalleled story

- 1999 Graphics computation pipeline implemented in dedicated *graphics hardware*
 - Computations using OpenGL graphics library [2]
 - »GPU« coined by NVIDIA [3]
- 2001 NVIDIA GeForce 3 with *programmable* shaders (instead of fixed pipeline) and floating-point support; 2003: DirectX 9 at ATI

History of GPUs

A short but unparalleled story

- 1999 Graphics computation pipeline implemented in dedicated *graphics hardware*
 - Computations using OpenGL graphics library [2]
 - »GPU« coined by NVIDIA [3]
- 2001 NVIDIA GeForce 3 with *programmable* shaders (instead of fixed pipeline) and floating-point support; 2003: DirectX 9 at ATI
- 2007 CUDA

History of GPUs

A short but unparalleled story

- 1999 Graphics computation pipeline implemented in dedicated *graphics hardware*
 - Computations using OpenGL graphics library [2]
 - »GPU« coined by NVIDIA [3]
- 2001 NVIDIA GeForce 3 with *programmable* shaders (instead of fixed pipeline) and floating-point support; 2003: DirectX 9 at ATI
- 2007 CUDA
- 2009 OpenCL

History of GPUs

A short but unparalleled story

- 1999 Graphics computation pipeline implemented in dedicated *graphics hardware*
 - Computations using OpenGL graphics library [2]
 - »GPU« coined by NVIDIA [3]
- 2001 NVIDIA GeForce 3 with *programmable* shaders (instead of fixed pipeline) and floating-point support; 2003: DirectX 9 at ATI
- 2007 CUDA
- 2009 OpenCL
- 2022 Top 500: 32 % with GPUs (#1, #2; 7 of top 10) [4], Green 500: 9 of top 10 with GPUs [5]

History of GPUs

A short but unparalleled story

- 1999 Graphics computation pipeline implemented in dedicated *graphics hardware*
 - Computations using OpenGL graphics library [2]
 - »GPU« coined by NVIDIA [3]
- 2001 NVIDIA GeForce 3 with *programmable* shaders (instead of fixed pipeline) and floating-point support; 2003: DirectX 9 at ATI
- 2007 CUDA
- 2009 OpenCL
- 2022 Top 500: 32 % with GPUs (#1, #2; 7 of top 10) [4], Green 500: 9 of top 10 with GPUs [5]
- 2023 🇮🇹: Leonardo (238 PFLOP/s*, Italy), NVIDIA GPUs; LUMI (309 PFLOP/s*, Finland), AMD GPUs
🇺🇸: Frontier (1.102 EFLOP/s*, ORNL), AMD GPUs

*: Effective FLOP/s, not theoretical peak (HPL R_{max})

History of GPUs

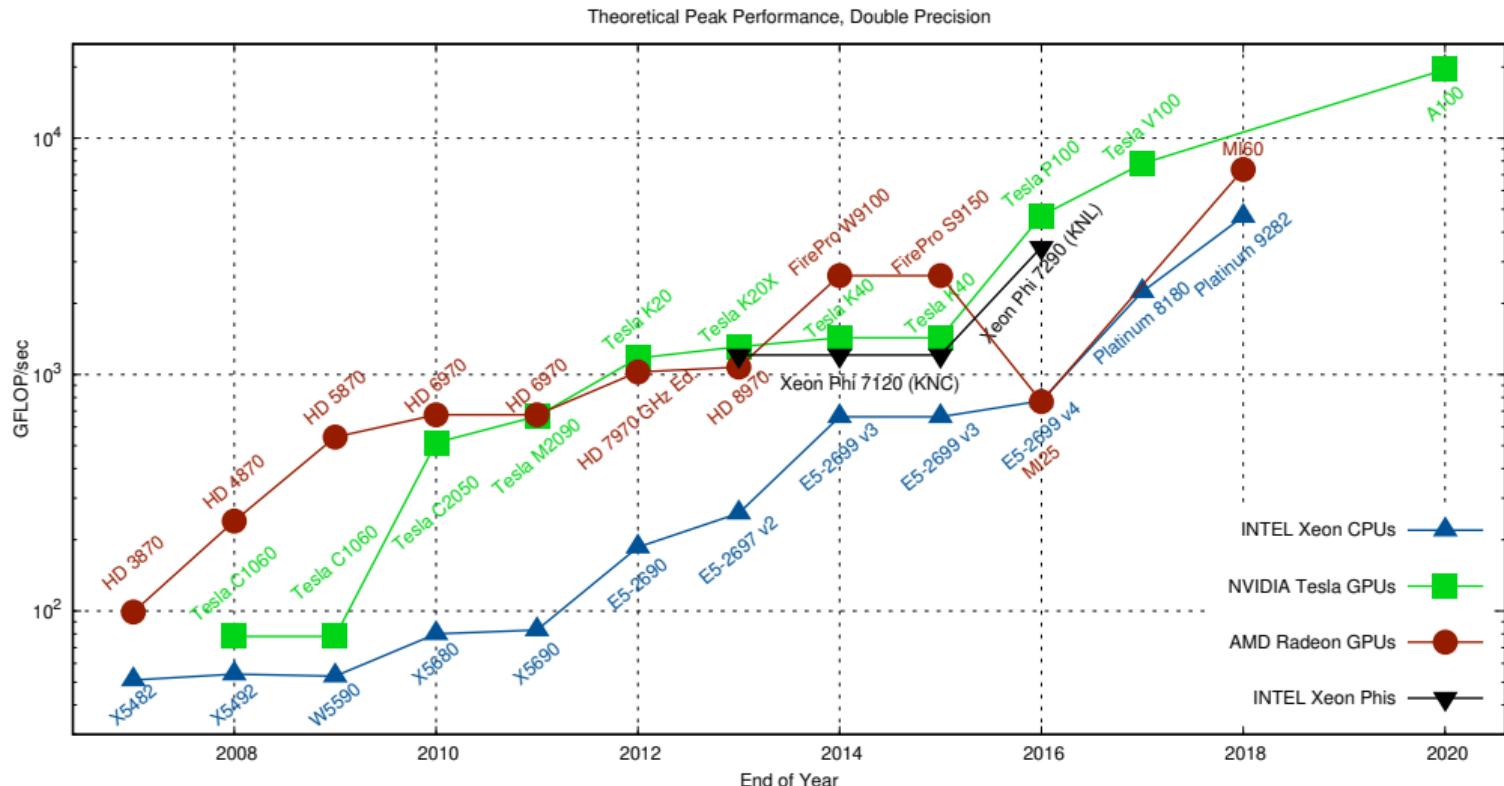
A short but unparalleled story

- 1999 Graphics computation pipeline implemented in dedicated *graphics hardware*
 - Computations using OpenGL graphics library [2]
 - »GPU« coined by NVIDIA [3]
- 2001 NVIDIA GeForce 3 with *programmable* shaders (instead of fixed pipeline) and floating-point support; 2003: DirectX 9 at ATI
- 2007 CUDA
- 2009 OpenCL
- 2022 Top 500: 32 % with GPUs (#1, #2; 7 of top 10) [4], Green 500: 9 of top 10 with GPUs [5]
- 2023 : Leonardo (238 PFLOP/s*, Italy), NVIDIA GPUs; LUMI (309 PFLOP/s*, Finland), AMD GPUs
: Frontier (1.102 EFLOP/s*, ORNL), AMD GPUs
- Soon : JUPITER (\approx 1 EFLOP/s, NVIDIA GPUs, JSC)
: Aurora (\approx 2 EFLOP/s, Argonne), Intel GPUs; El Capitan (\approx 2 EFLOP/s, LLNL), AMD GPUs

*: Effective FLOP/s, not theoretical peak (HPL R_{max})

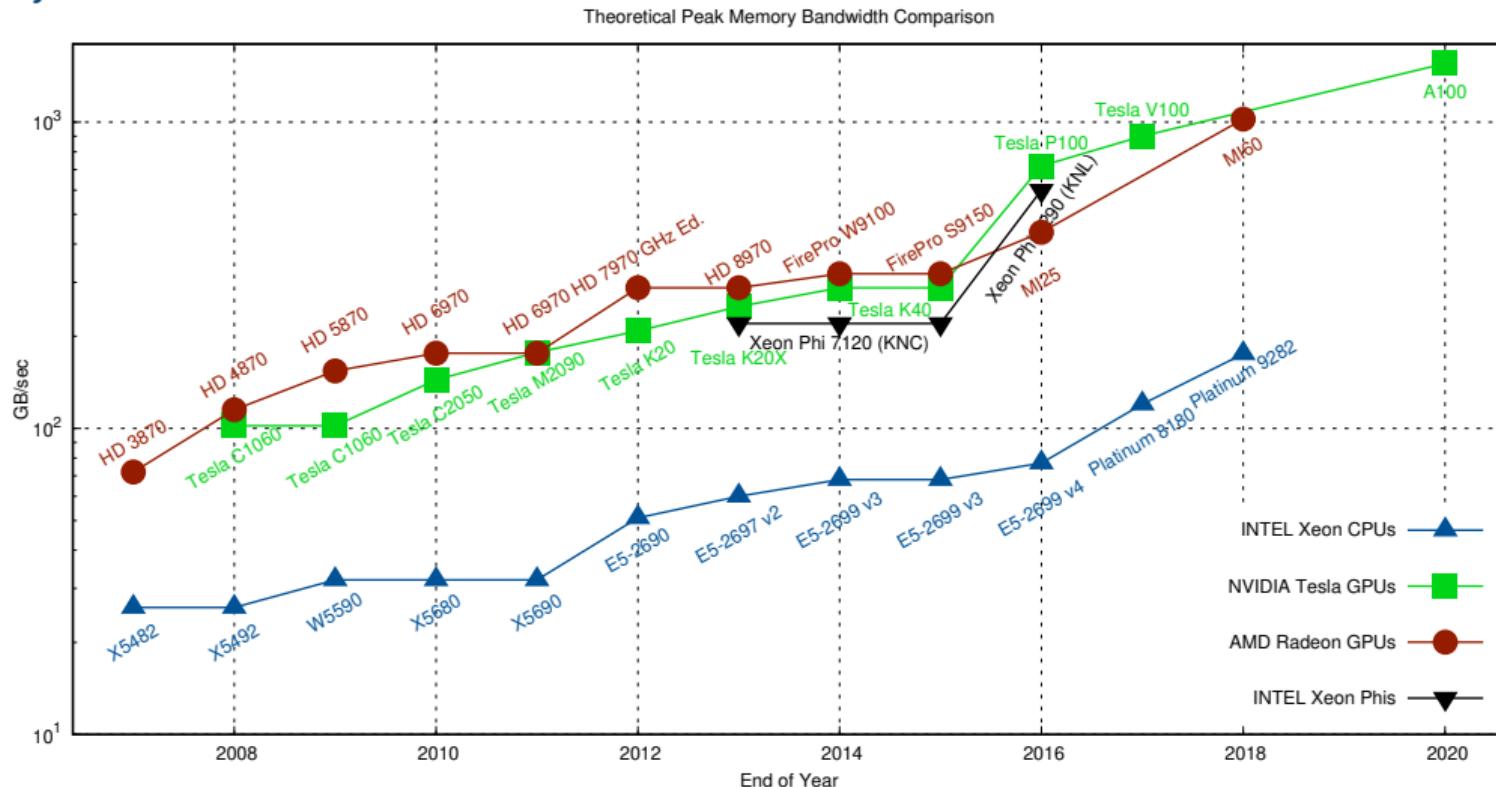
Status Quo Across Architectures

Performance



Status Quo Across Architectures

Memory Bandwidth



Graphic: Rupp [6]



JUWELS Cluster – Jülich's Scalable System

- 2500 nodes with Intel Xeon CPUs (2×24 cores)
- 46 + 10 nodes with 4 NVIDIA Tesla V100 cards (16 GB memory)
- 10.4 (CPU) + 1.6 (GPU) PFLOP/s peak performance (Top500: #86)



JUWELS Booster – Scaling Higher!

- 936 nodes with AMD EPYC Rome CPUs (2×24 cores)
- Each with 4 NVIDIA A100 Ampere GPUs (each: $\text{FP64TC: } 19.5$ TFLOP/s, 40 GB memory)
 $\text{FP64: } 9.7$
- InfiniBand DragonFly+ HDR-200 network; 4×200 Gbit/s per node

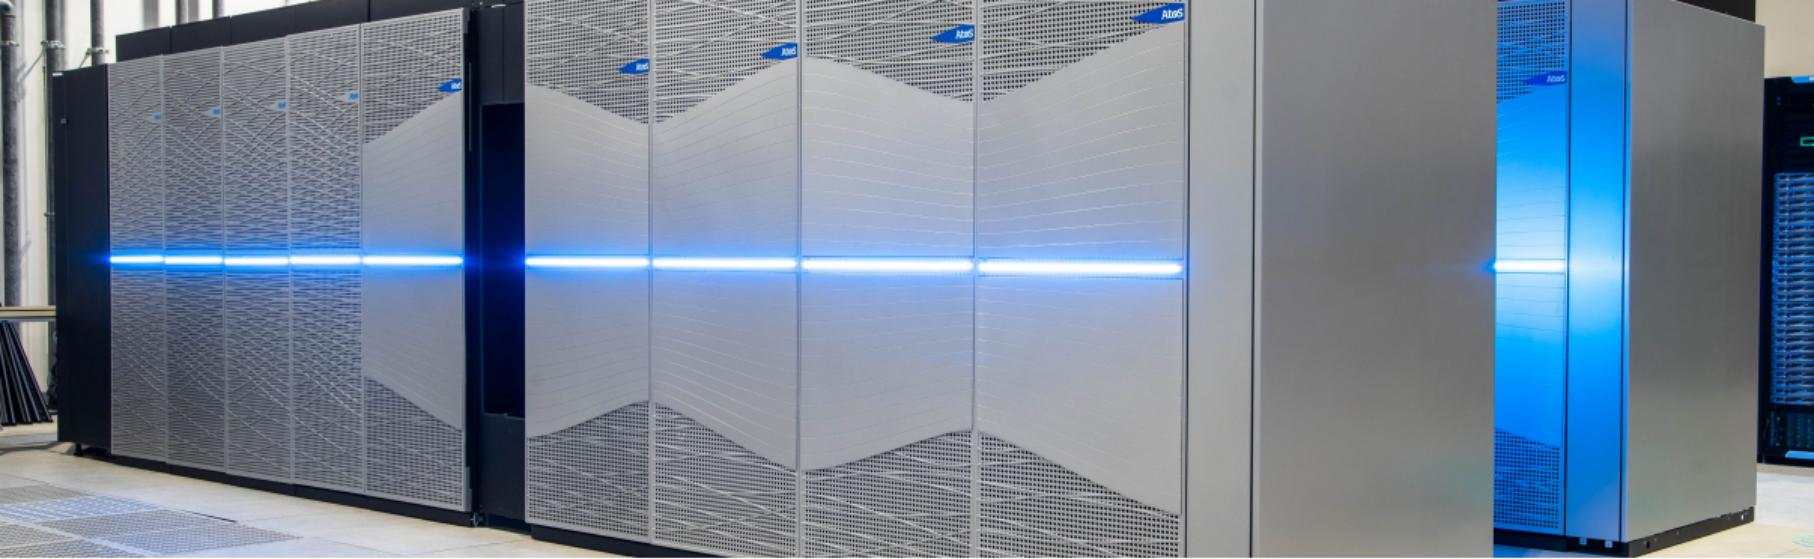


Top500 List Nov 2020:

- #1 Europe
- #7 World
- #4* Top/Green500

JUWELS Booster – Scaling Higher!

- 936 nodes with AMD EPYC Rome CPUs (2×24 cores)
- Each with 4 NVIDIA A100 Ampere GPUs (each: $\text{FP64TC: } 19.5$ TFLOP/s, 40 GB memory)
 $\text{FP64: } 9.7$
- InfiniBand DragonFly+ HDR-200 network; 4×200 Gbit/s per node



JURECA DC – Multi-Purpose

- 768 nodes with AMD EPYC Rome CPUs (2×64 cores)
- 192 nodes with 4 NVIDIA A100 Ampere GPUs
- InfiniBand DragonFly+ HDR-100 network

Getting GPU-Acquainted

Some Applications

TASK

Location of Code:

1-Introduction-GPU-Programming/Tasks/getting-started

See Instructions.ipynb for hints.

Make sure to have sourced the course environment!

Getting GPU-Acquainted

TASK

Some Applications

GEMM

N-Body

Location of Code:

1-Introduction-GPU-Programming/Tasks/getting-started

See Instructions.ipynb for hints.

Make sure to have sourced the course environment!

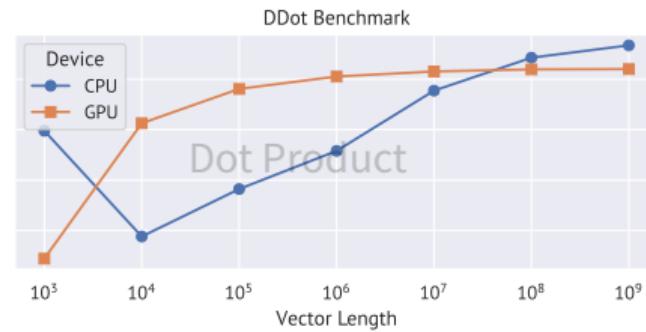
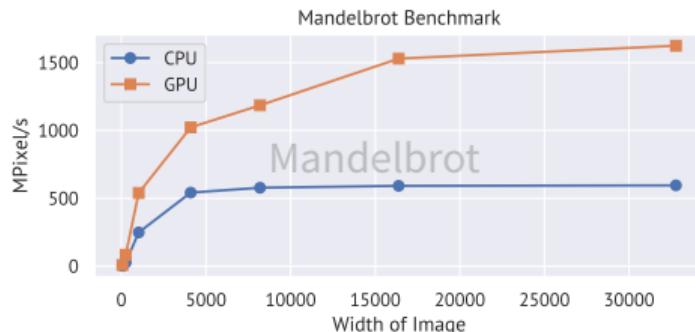
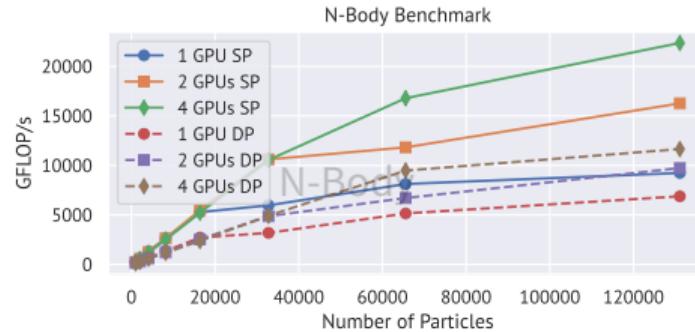
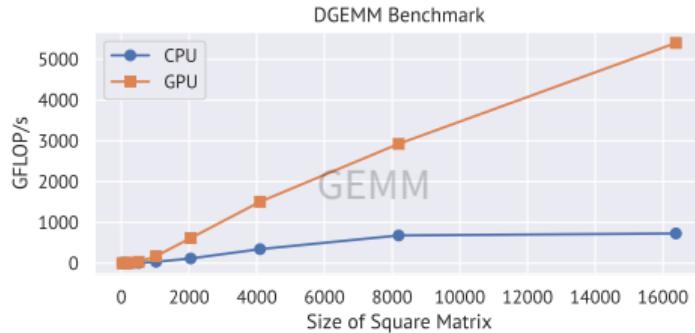
Mandelbrot

Dot Product

Getting GPU-Acquainted

TASK

Some Applications



Platform

CPU vs. GPU

A matter of specialties



CPU vs. GPU

A matter of specialties



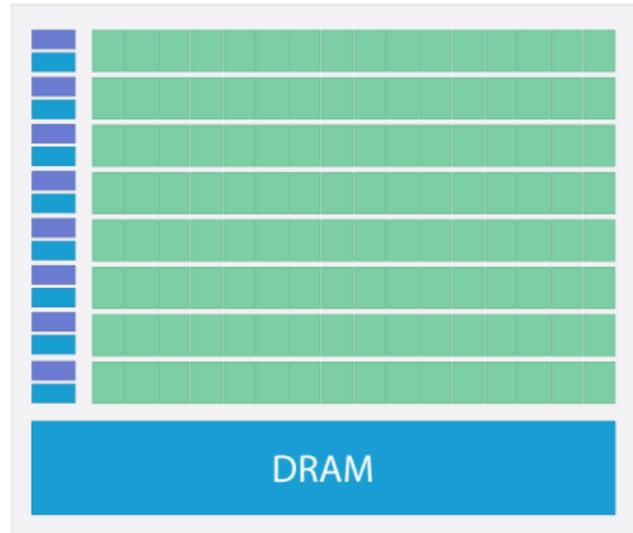
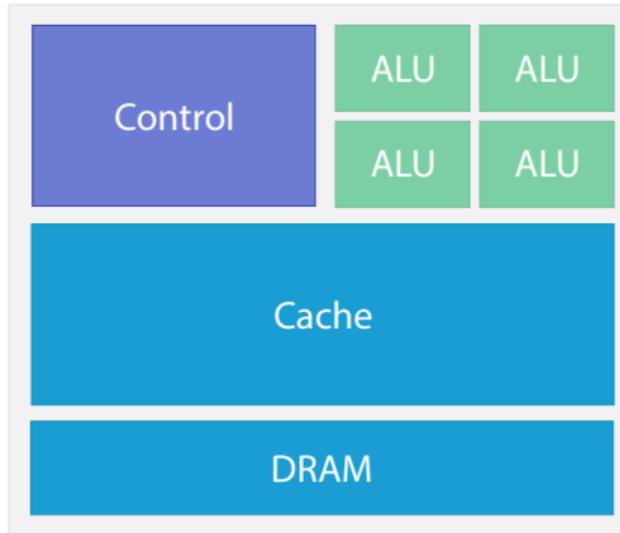
Transporting one



Transporting many

CPU vs. GPU

Chip



GPU Architecture

Overview

Aim: Hide Latency
Everything else follows

GPU Architecture

Overview

Aim: Hide Latency
Everything else follows

SIMT

Asynchronicity

Memory

GPU Architecture

Overview

Aim: Hide Latency
Everything else follows

SIMT

Asynchronicity

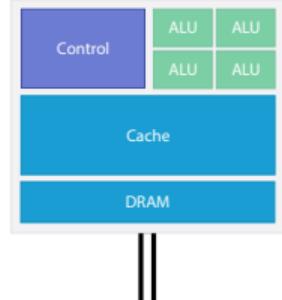
Memory

Memory

GPU memory ain't no CPU memory

- GPU: accelerator / extension card
- Separate device from CPU

Host



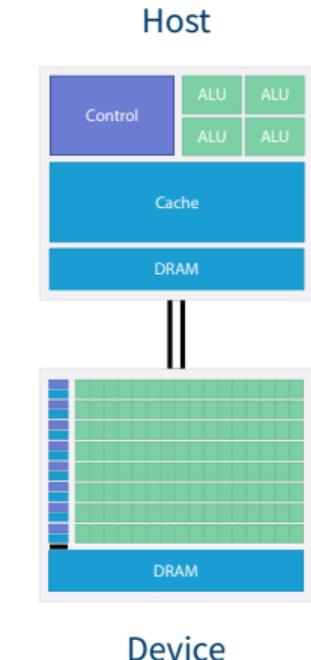
Device

Memory

GPU memory ain't no CPU memory

Unified Virtual Addressing

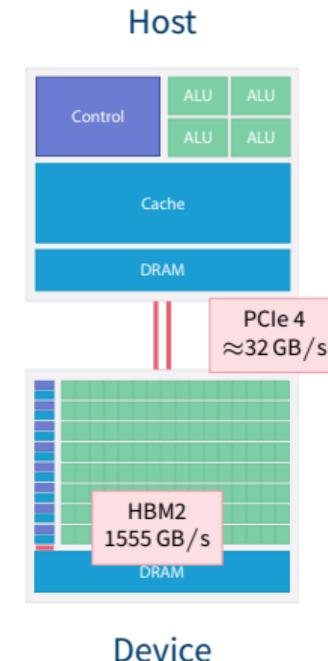
- GPU: accelerator / extension card
 - Separate device from CPU
- Separate memory, but UVA**



Memory

GPU memory ain't no CPU memory

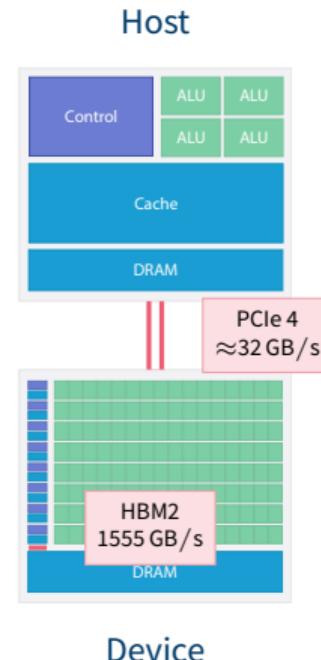
- GPU: accelerator / extension card
- Separate device from CPU
- Separate memory, but UVA**



Memory

GPU memory ain't no CPU memory

- GPU: accelerator / extension card
- Separate device from CPU
- Separate memory, but UVA**
- Memory transfers need special consideration!
Do as little as possible!

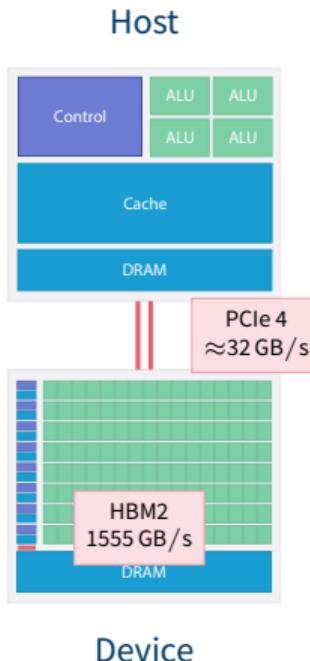


Memory

GPU memory ain't no CPU memory

Unified Memory

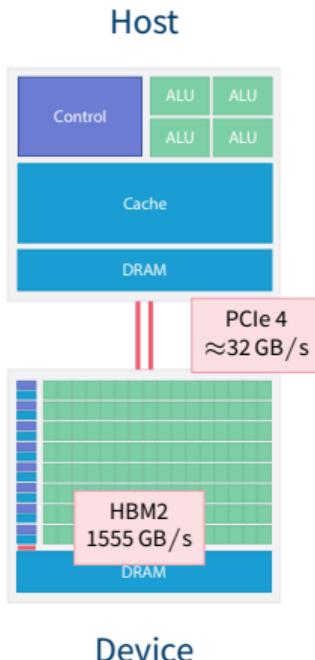
- GPU: accelerator / extension card
- Separate device from CPU
- Separate memory, but UVA and UM**
- Memory transfers need special consideration!
Do as little as possible!
- Choice: automatic transfers (convenience) or manual transfers (control)



Memory

GPU memory ain't no CPU memory

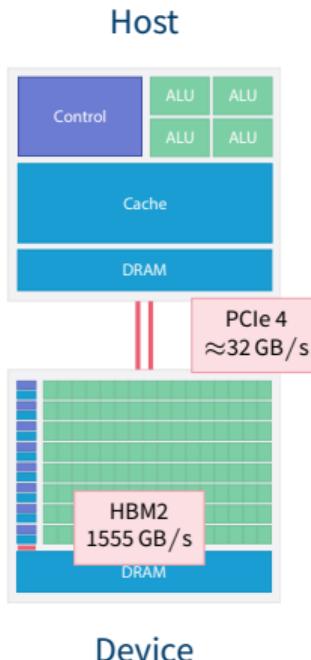
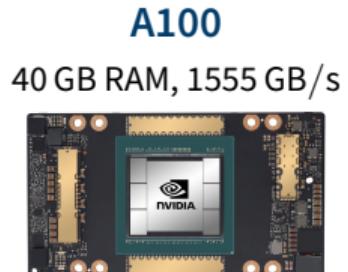
- GPU: accelerator / extension card
- Separate device from CPU
- Separate memory, but UVA and UM**
- Memory transfers need special consideration!
Do as little as possible!
- Choice: automatic transfers (convenience) or manual transfers (control)



Memory

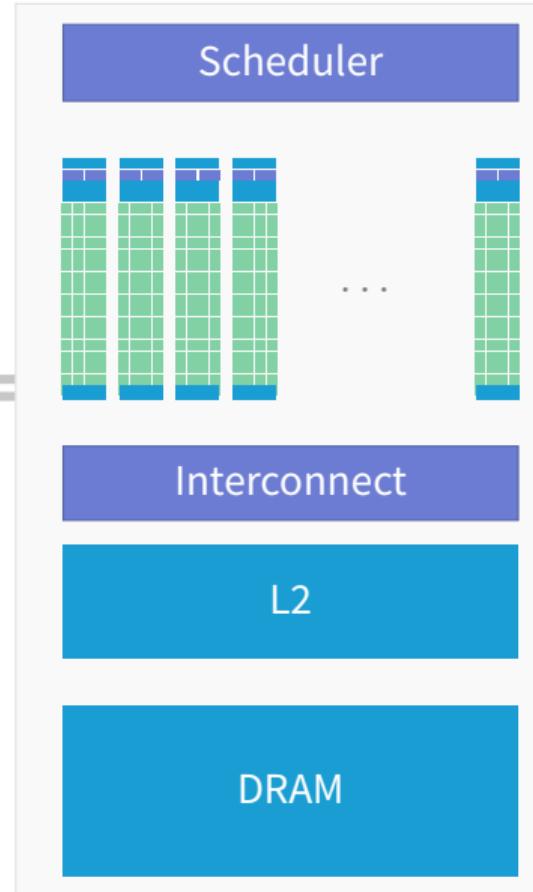
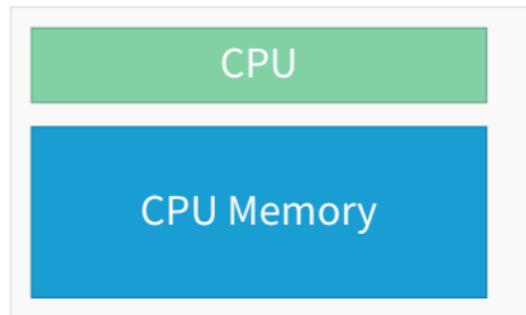
GPU memory ain't no CPU memory

- GPU: accelerator / extension card
- Separate device from CPU
- Separate memory, but UVA and UM**
- Memory transfers need special consideration!
Do as little as possible!
- Choice: automatic transfers (convenience) or manual transfers (control)



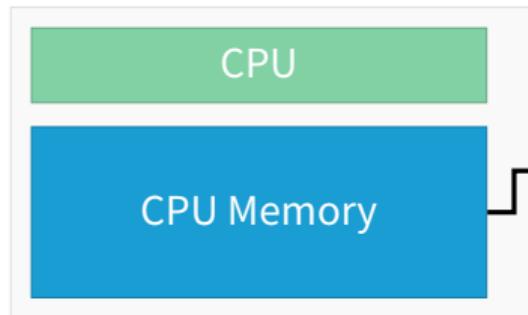
Processing Flow

CPU → GPU → CPU

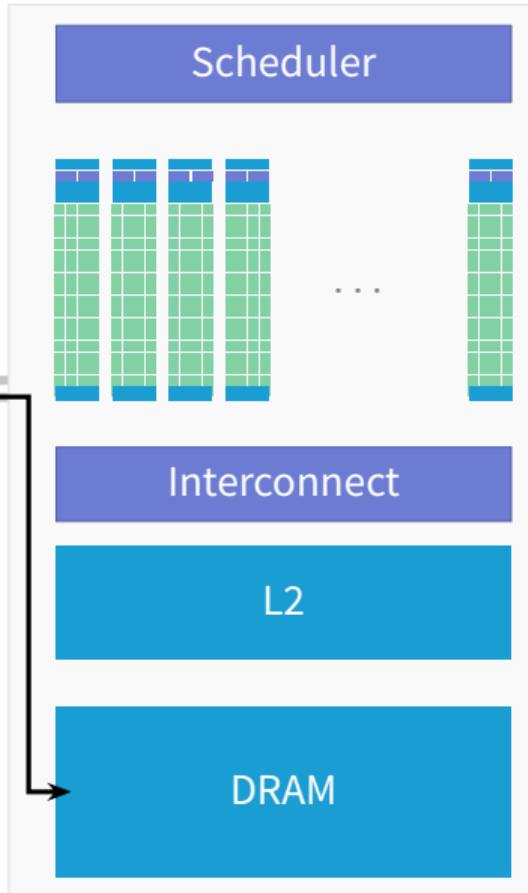


Processing Flow

CPU → GPU → CPU

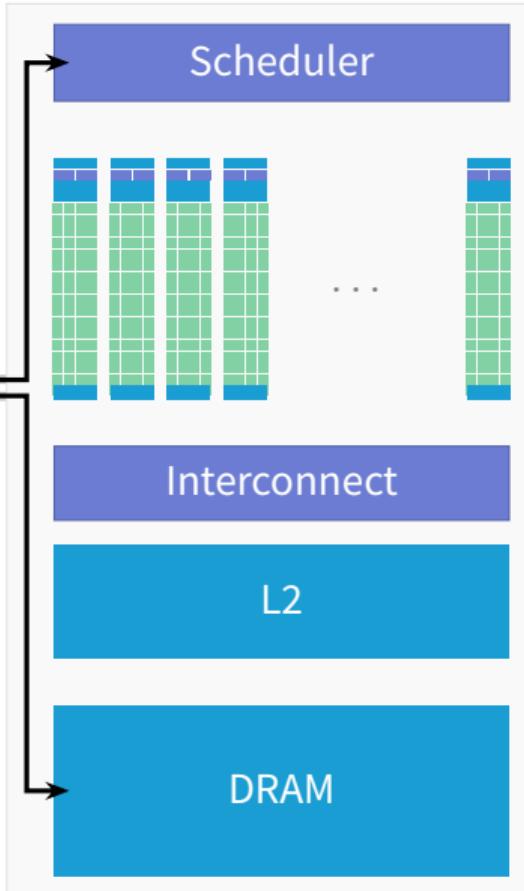
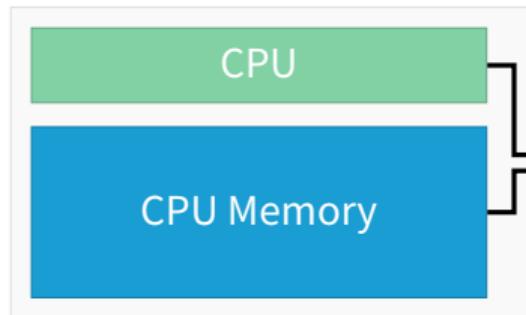


- 1 Transfer data from CPU memory to GPU memory



Processing Flow

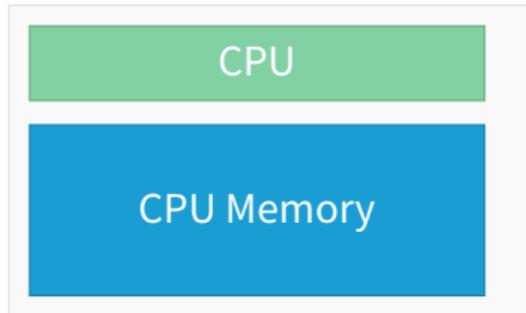
CPU → GPU → CPU



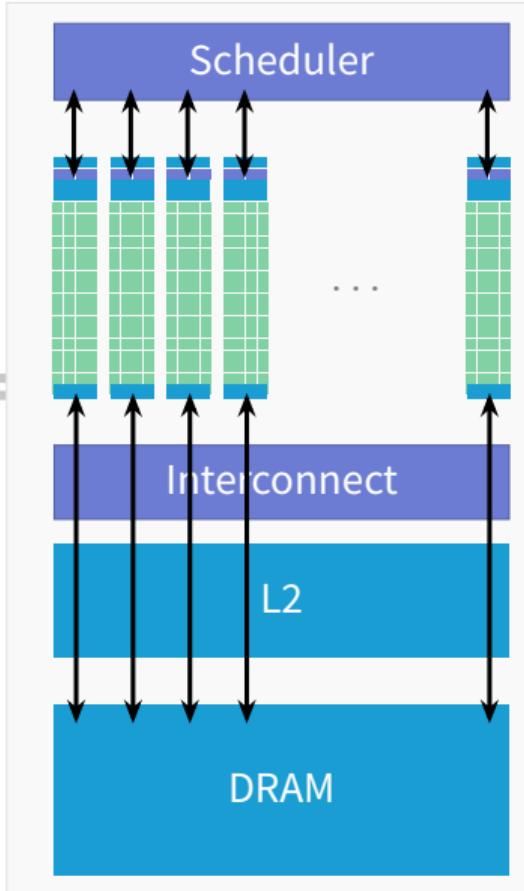
- 1 Transfer data from CPU memory to GPU memory, transfer program

Processing Flow

CPU → GPU → CPU

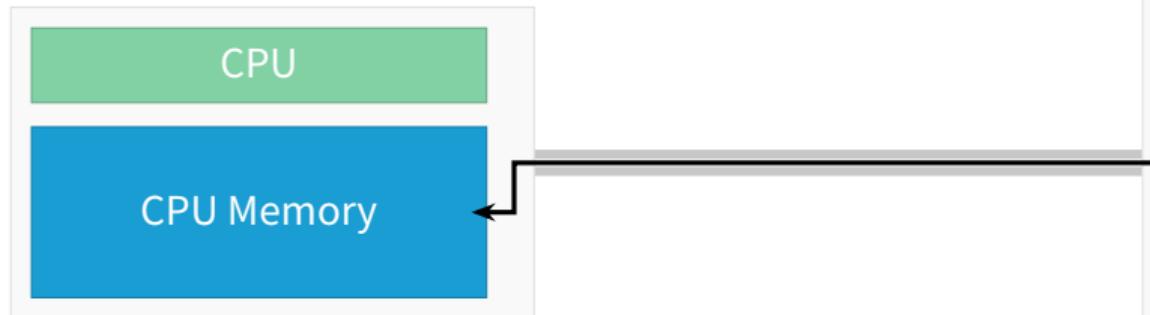


- 1 Transfer data from CPU memory to GPU memory, transfer program
- 2 Load GPU program, execute on SMs, get (cached) data from memory; write back

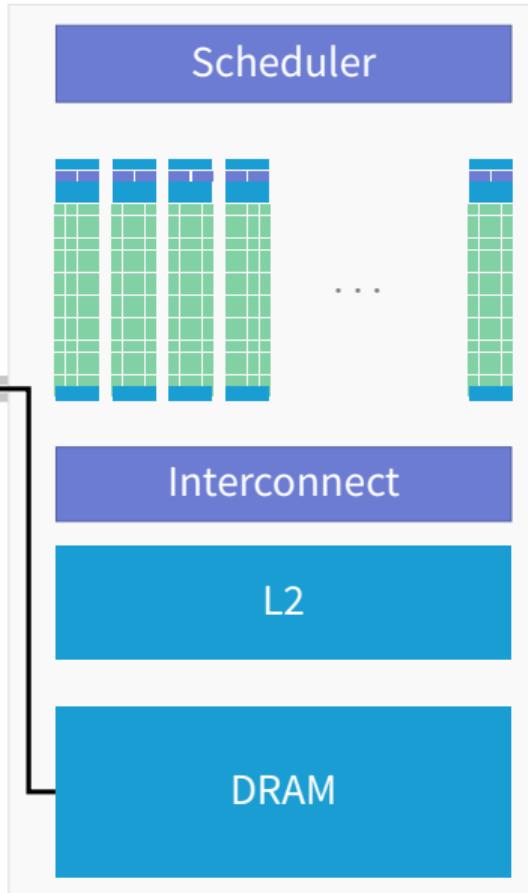


Processing Flow

CPU → GPU → CPU



- 1 Transfer data from CPU memory to GPU memory, transfer program
- 2 Load GPU program, execute on SMs, get (cached) data from memory; write back
- 3 Transfer results back to host memory



GPU Architecture

Overview

Aim: Hide Latency
Everything else follows

SIMT

Asynchronicity

Memory

GPU Architecture

Overview

Aim: Hide Latency
Everything else follows

SIMT

Asynchronicity

Memory

Async

Following different streams

- Problem: Memory transfer is comparably slow
Solution: Do something else in meantime (**computation**)!
- Overlap tasks
- Copy and compute engines run separately (*streams*)



- GPU needs to be fed: Schedule many computations
- CPU can do other work while GPU computes; synchronization

GPU Architecture

Overview

Aim: Hide Latency
Everything else follows

SIMT

Asynchronicity

Memory

GPU Architecture

Overview

Aim: Hide Latency
Everything else follows

SIMT

Asynchronicity

Memory

Flynn's Taxonomy

- Michael Flynn (1966/1972): classification of computer architectures
- Define by number of instructions operating on data elements

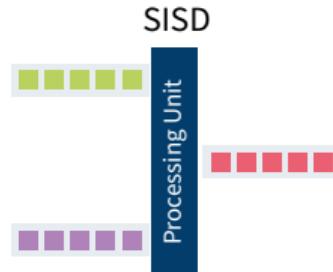
Flynn's Taxonomy

- Michael Flynn (1966/1972): classification of computer architectures
- Define by number of instructions operating on data elements
- $\begin{pmatrix} \text{Single} \\ \text{Multiple} \end{pmatrix} \otimes \begin{pmatrix} \text{Instruction} \\ \text{Data} \end{pmatrix}$

Flynn's Taxonomy

- Michael Flynn (1966/1972): classification of computer architectures
- Define by number of instructions operating on data elements
- $\begin{pmatrix} \text{Single} \\ \text{Multiple} \end{pmatrix} \otimes \begin{pmatrix} \text{Instruction} \\ \text{Data} \end{pmatrix}$
SISD Single Instruction, Single Data

Flynn's Taxonomy



- Michael Flynn (1966/1972): classification of computer architectures
- Define by number of instructions operating on data elements
- $\begin{pmatrix} \text{Single} \\ \text{Multiple} \end{pmatrix} \otimes \begin{pmatrix} \text{Instruction} \\ \text{Data} \end{pmatrix}$

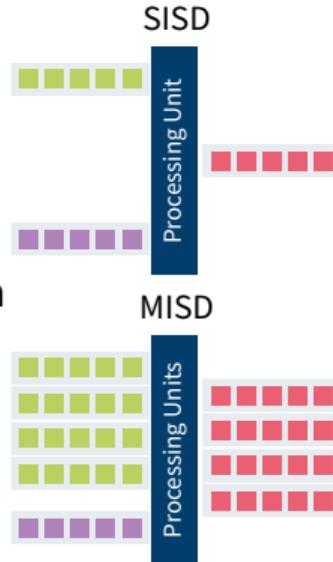
SISD Single Instruction, Single Data

Flynn's Taxonomy

- Michael Flynn (1966/1972): classification of computer architectures
- Define by number of instructions operating on data elements
- $\begin{pmatrix} \text{Single} \\ \text{Multiple} \end{pmatrix} \otimes \begin{pmatrix} \text{Instruction} \\ \text{Data} \end{pmatrix}$

SISD Single Instruction, Single Data

MISD Multiple Instructions, Single Data



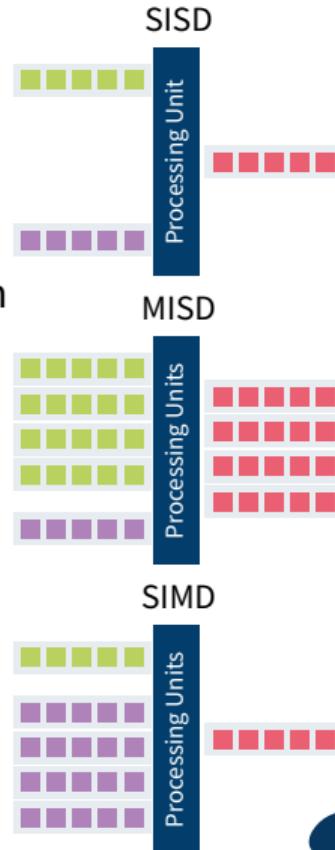
Flynn's Taxonomy

- Michael Flynn (1966/1972): classification of computer architectures
- Define by number of instructions operating on data elements
- $\begin{pmatrix} \text{Single} \\ \text{Multiple} \end{pmatrix} \otimes \begin{pmatrix} \text{Instruction} \\ \text{Data} \end{pmatrix}$

SISD Single Instruction, Single Data

MISD Multiple Instructions, Single Data

SIMD Single Instruction, Multiple Data



Flynn's Taxonomy

- Michael Flynn (1966/1972): classification of computer architectures
- Define by number of instructions operating on data elements
- $(\begin{matrix} \text{Single} \\ \text{Multiple} \end{matrix}) \otimes (\begin{matrix} \text{Instruction} \\ \text{Data} \end{matrix})$

SISD Single Instruction, Single Data

MISD Multiple Instructions, Single Data

SIMD Single Instruction, Multiple Data

MIMD Multiple Instructions, Multiple Data



Flynn's Taxonomy

- Michael Flynn (1966/1972): classification of computer architectures
- Define by number of instructions operating on data elements
- $\begin{pmatrix} \text{Single} \\ \text{Multiple} \end{pmatrix} \otimes \begin{pmatrix} \text{Instruction} \\ \text{Data} \end{pmatrix}$

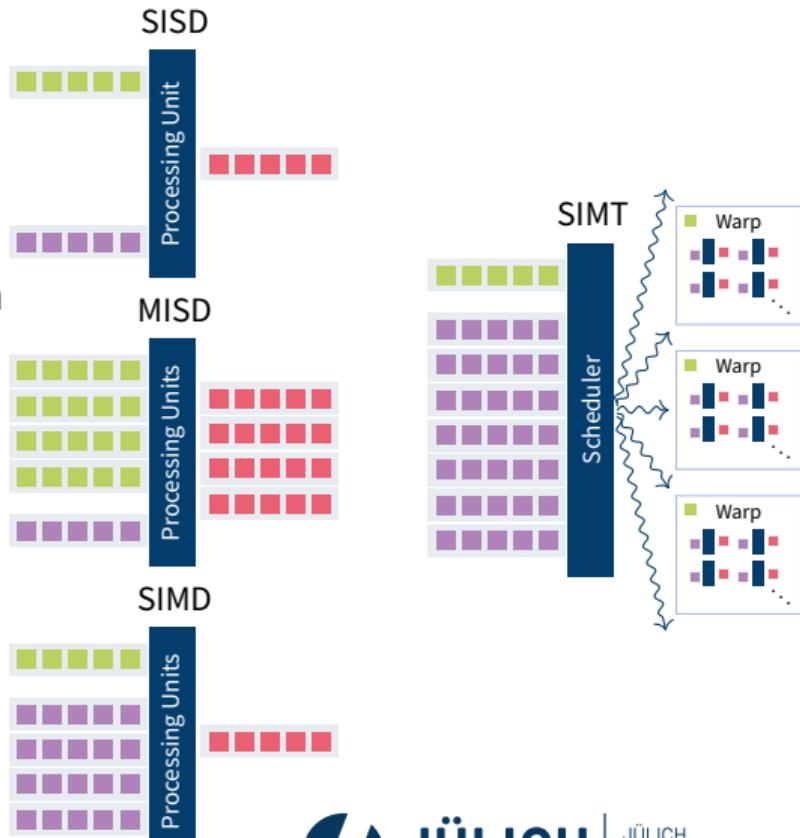
SISD Single Instruction, Single Data

MISD Multiple Instructions, Single Data

SIMD Single Instruction, Multiple Data

MIMD Multiple Instructions, Multiple Data

SIMT Single Instruction, Multiple *Threads*



SIMT

SIMT = SIMD \oplus SMT

- CPU:
 - Single Instruction, Multiple Data (SIMD)

Scalar

$$\begin{array}{rcl} A_0 & + & B_0 \\ \hline A_1 & + & B_1 \\ \hline A_2 & + & B_2 \\ \hline A_3 & + & B_3 \\ \hline \end{array} = \begin{array}{l} C_0 \\ C_1 \\ C_2 \\ C_3 \end{array}$$

SIMT

SIMT = SIMD \oplus SMT

- CPU:
 - Single Instruction, Multiple Data (SIMD)

Vector

$$\begin{array}{c} A_0 \\ A_1 \\ A_2 \\ A_3 \end{array} + \begin{array}{c} B_0 \\ B_1 \\ B_2 \\ B_3 \end{array} = \begin{array}{c} C_0 \\ C_1 \\ C_2 \\ C_3 \end{array}$$

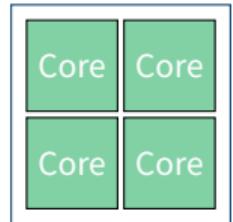
SIMT

SIMT = SIMD \oplus SMT

- CPU:
 - Single Instruction, Multiple Data (SIMD)
 - Simultaneous Multithreading (SMT)

Vector

$$\begin{array}{c} A_0 \\ A_1 \\ A_2 \\ A_3 \end{array} + \begin{array}{c} B_0 \\ B_1 \\ B_2 \\ B_3 \end{array} = \begin{array}{c} C_0 \\ C_1 \\ C_2 \\ C_3 \end{array}$$



SIMT

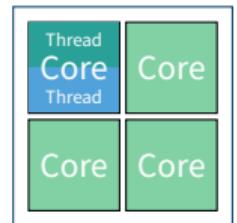
SIMT = SIMD \oplus SMT

- CPU:
 - Single Instruction, Multiple Data (SIMD)
 - Simultaneous Multithreading (SMT)

Vector

$$\begin{array}{c} A_0 \\ A_1 \\ A_2 \\ A_3 \end{array} + \begin{array}{c} B_0 \\ B_1 \\ B_2 \\ B_3 \end{array} = \begin{array}{c} C_0 \\ C_1 \\ C_2 \\ C_3 \end{array}$$

SMT



SIMT

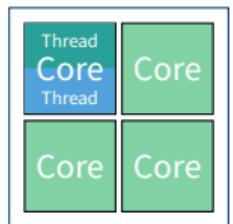
SIMT = SIMD \oplus SMT

- CPU:
 - Single Instruction, Multiple Data (SIMD)
 - Simultaneous Multithreading (SMT)
- GPU: Single Instruction, Multiple Threads (SIMT)

Vector

$$\begin{array}{c} A_0 \\ A_1 \\ A_2 \\ A_3 \end{array} + \begin{array}{c} B_0 \\ B_1 \\ B_2 \\ B_3 \end{array} = \begin{array}{c} C_0 \\ C_1 \\ C_2 \\ C_3 \end{array}$$

SMT



SIMT

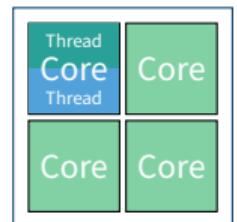
SIMT = SIMD \oplus SMT

- CPU:
 - Single Instruction, Multiple Data (SIMD)
 - Simultaneous Multithreading (SMT)
- GPU: Single Instruction, Multiple Threads (SIMT)

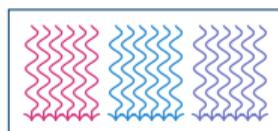
Vector

$$\begin{array}{c} A_0 \\ A_1 \\ A_2 \\ A_3 \end{array} + \begin{array}{c} B_0 \\ B_1 \\ B_2 \\ B_3 \end{array} = \begin{array}{c} C_0 \\ C_1 \\ C_2 \\ C_3 \end{array}$$

SMT



SIMT

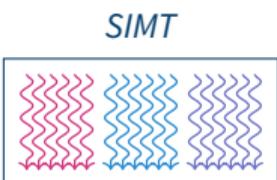
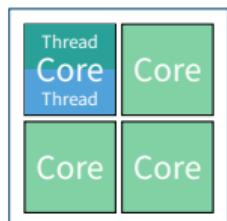


SIMT

$$\text{SIMT} = \text{SIMD} \oplus \text{SMT}$$

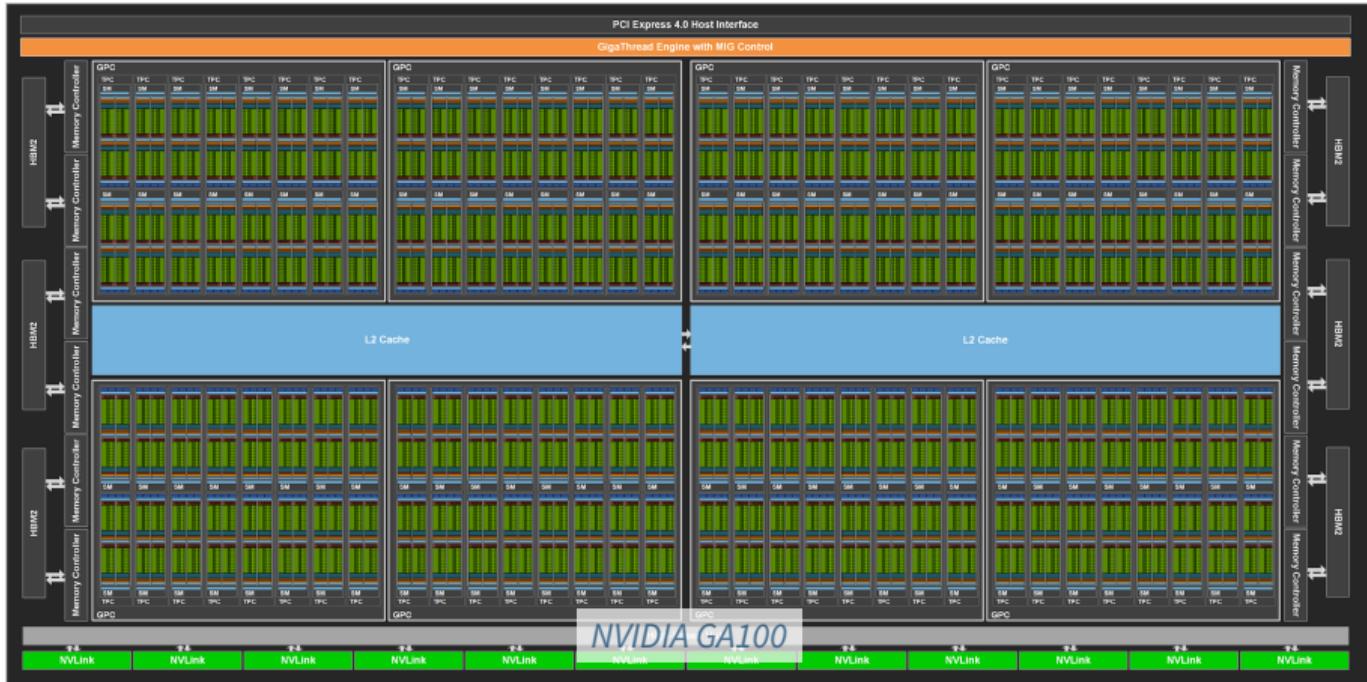
- CPU:
 - Single Instruction, Multiple Data (**SIMD**)
 - Simultaneous Multithreading (**SMT**)
 - GPU: Single Instruction, Multiple Threads (**SIMT**)
 - CPU core \approx GPU multiprocessor (**SM**)
 - Working unit: set of threads (32, a *warp*)
 - Fast switching of threads (large register file)
 - Branching 

$$\begin{matrix} A_0 & B_0 & C_0 \\ A_1 & B_1 & C_1 \\ A_2 & B_2 & C_2 \\ A_3 & B_3 & C_3 \end{matrix}$$



SIMT

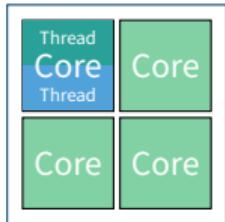
SIMT = SIMD \oplus SMT



Vector

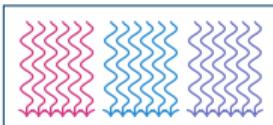
$$\begin{matrix} A_0 \\ A_1 \\ A_2 \\ A_3 \end{matrix} + \begin{matrix} B_0 \\ B_1 \\ B_2 \\ B_3 \end{matrix} = \begin{matrix} C_0 \\ C_1 \\ C_2 \\ C_3 \end{matrix}$$

SMT



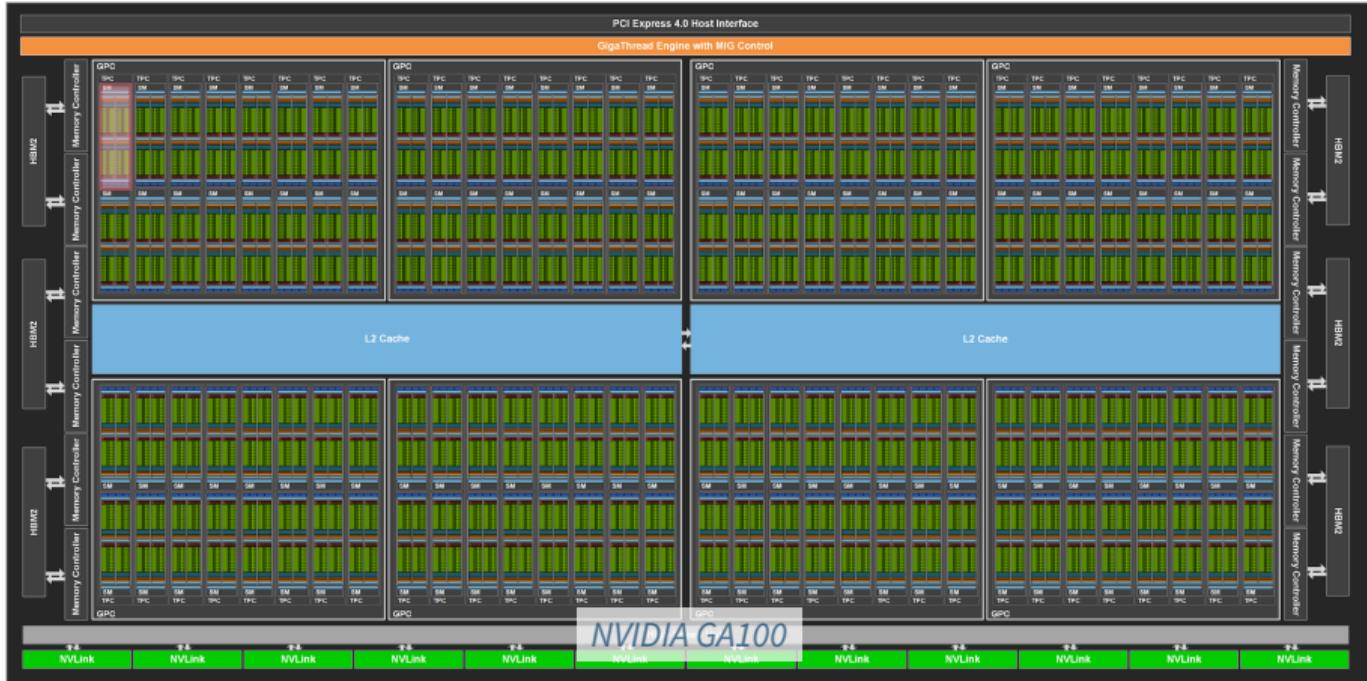
Graphics: img:ampere/pictures

SIMT



SIMT

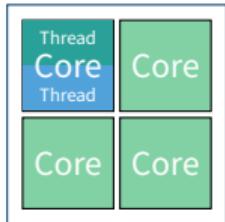
SIMT = SIMD \oplus SMT



Vector

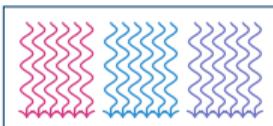
$$\begin{matrix} A_0 \\ A_1 \\ A_2 \\ A_3 \end{matrix} + \begin{matrix} B_0 \\ B_1 \\ B_2 \\ B_3 \end{matrix} = \begin{matrix} C_0 \\ C_1 \\ C_2 \\ C_3 \end{matrix}$$

SMT



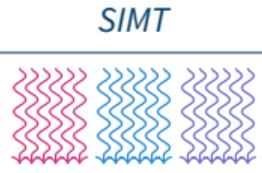
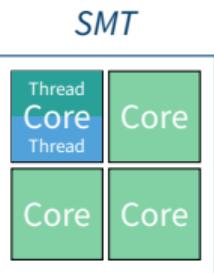
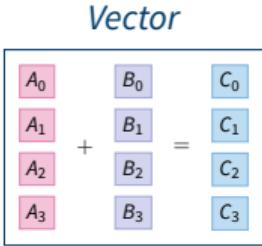
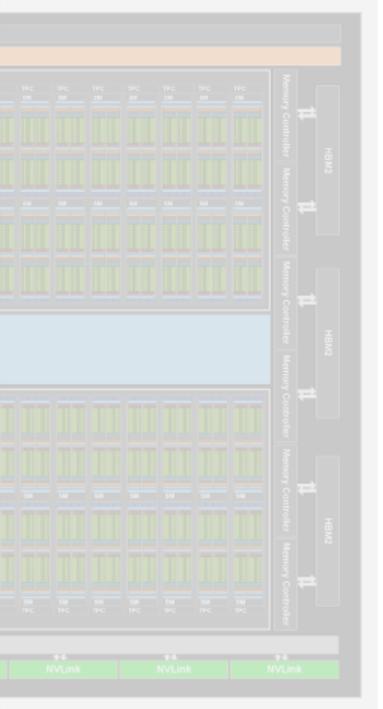
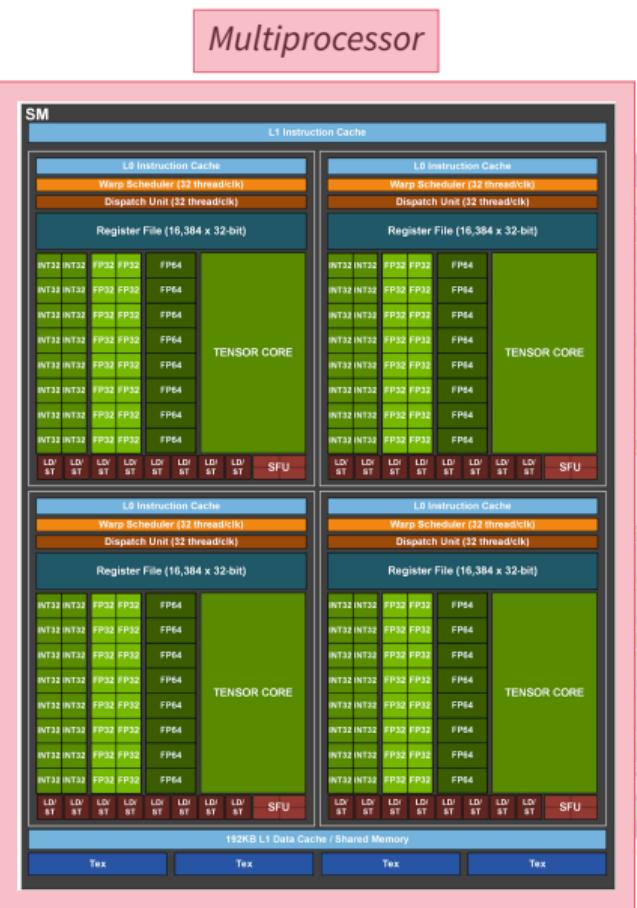
Graphics: img:ampere/pictures

SIMT



SIMT

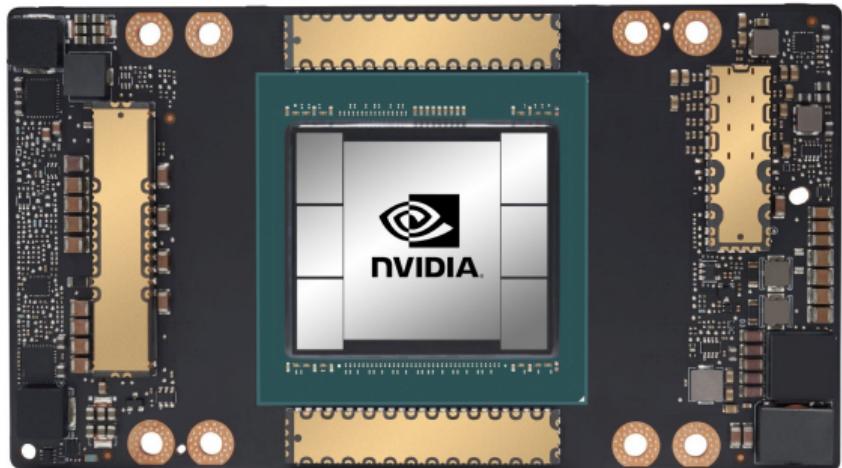
SIMT = SIMD \oplus SMT



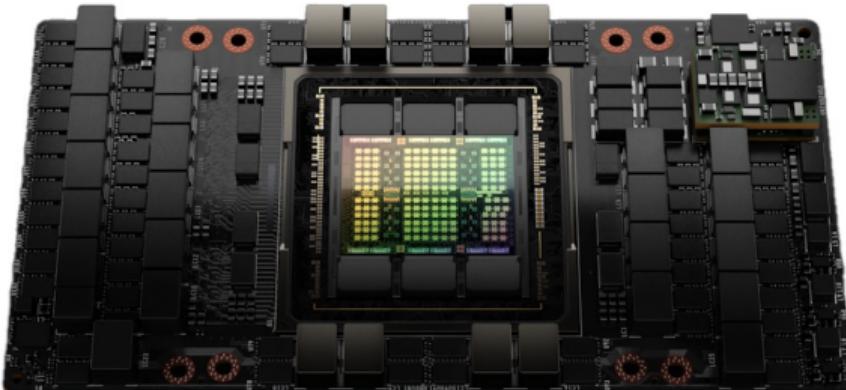
A100 vs H100

Comparison of current vs. next generation

A100

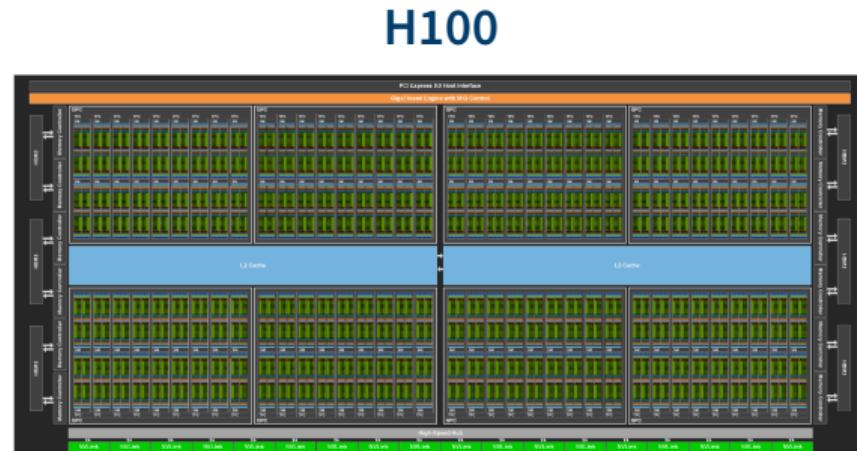


H100



A100 vs H100

Comparison of current vs. next generation



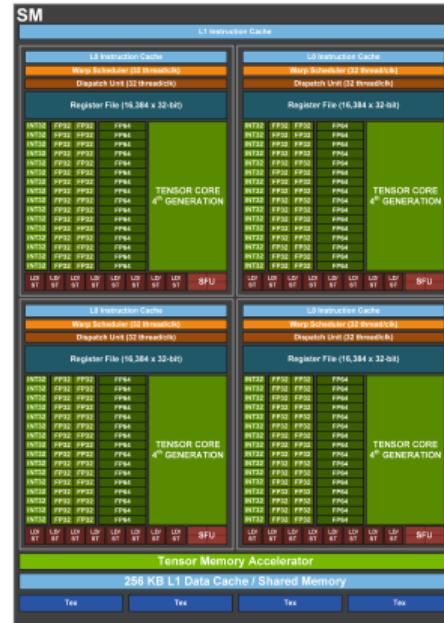
A100 vs H100

Comparison of current vs. next generation

A100



H100



Low Latency vs. High Throughput

Maybe GPU's ultimate feature

CPU Minimizes latency within each thread

GPU Hides latency with computations from other thread warps

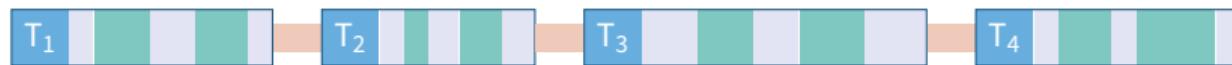
Low Latency vs. High Throughput

Maybe GPU's ultimate feature

CPU Minimizes latency within each thread

GPU Hides latency with computations from other thread warps

CPU Core: Low Latency



- Thread/Warp
- Processing
- Context Switch
- Ready
- Waiting

Low Latency vs. High Throughput

Maybe GPU's ultimate feature

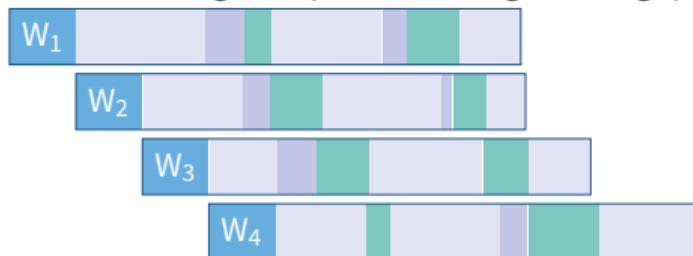
CPU Minimizes latency within each thread

GPU Hides latency with computations from other thread warps

CPU Core: Low Latency



GPU Streaming Multiprocessor: High Throughput



- █ Thread/Warp
- █ Processing
- █ Context Switch
- █ Ready
- █ Waiting

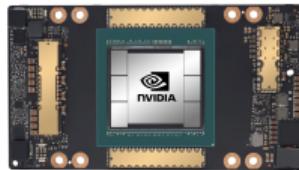
CPU vs. GPU

Let's summarize this!



Optimized for **low latency**

- + Large main memory
- + Fast clock rate
- + Large caches
- + Branch prediction
- + Powerful ALU
- Relatively low memory bandwidth
- Cache misses costly
- Low performance per watt



Optimized for **high throughput**

- + High bandwidth main memory
- + Latency tolerant (parallelism)
- + More compute resources
- + High performance per watt
- Limited memory capacity
- Low per-thread performance
- Extension card

Programming GPUs

Summary of Acceleration Possibilities

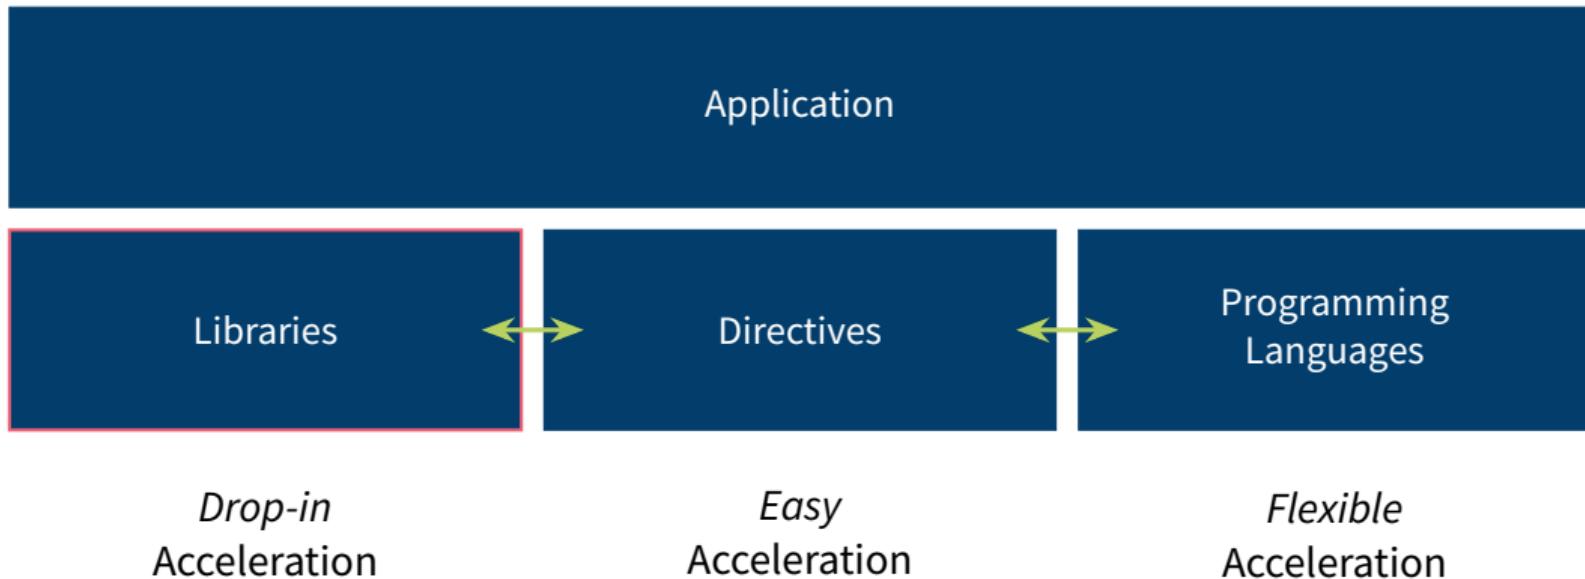


Drop-in
Acceleration

Easy
Acceleration

Flexible
Acceleration

Summary of Acceleration Possibilities



Libraries

Programming GPUs is easy: Just don't!

Libraries

Programming GPUs is easy: Just don't!

Use applications & libraries

Libraries

Programming GPUs is easy: Just don't!

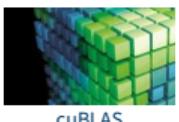
Use applications & libraries



Libraries

Programming GPUs is easy: Just don't!

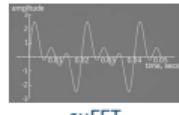
Use applications & libraries



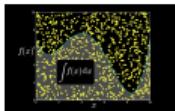
cuBLAS



cuSPARSE



cuFFT



cuRAND



CUDA Math



{ } ARRAYFIRE

Numba



CuPy

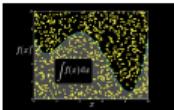
Libraries

Programming GPUs is easy: Just don't!

Use applications & libraries



cuFFT



cuRAND



{ } ARRAYFIRE



CUDA Math



Numba

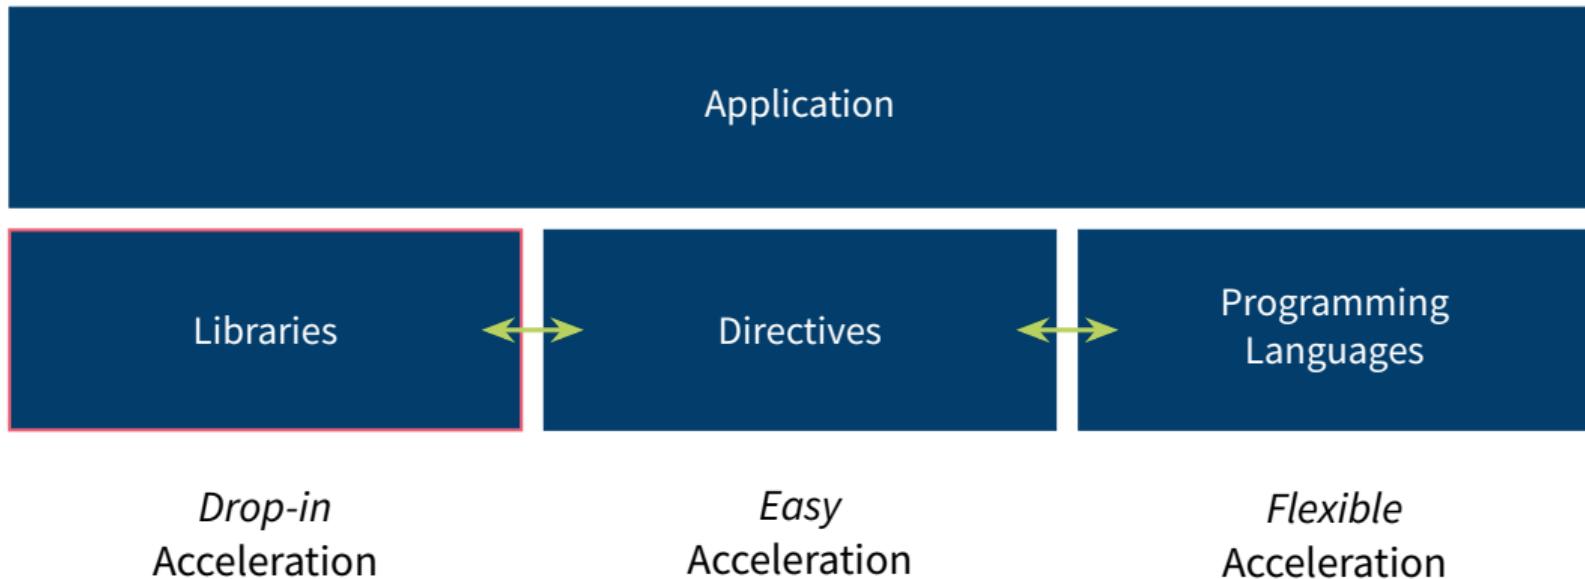


CuPy

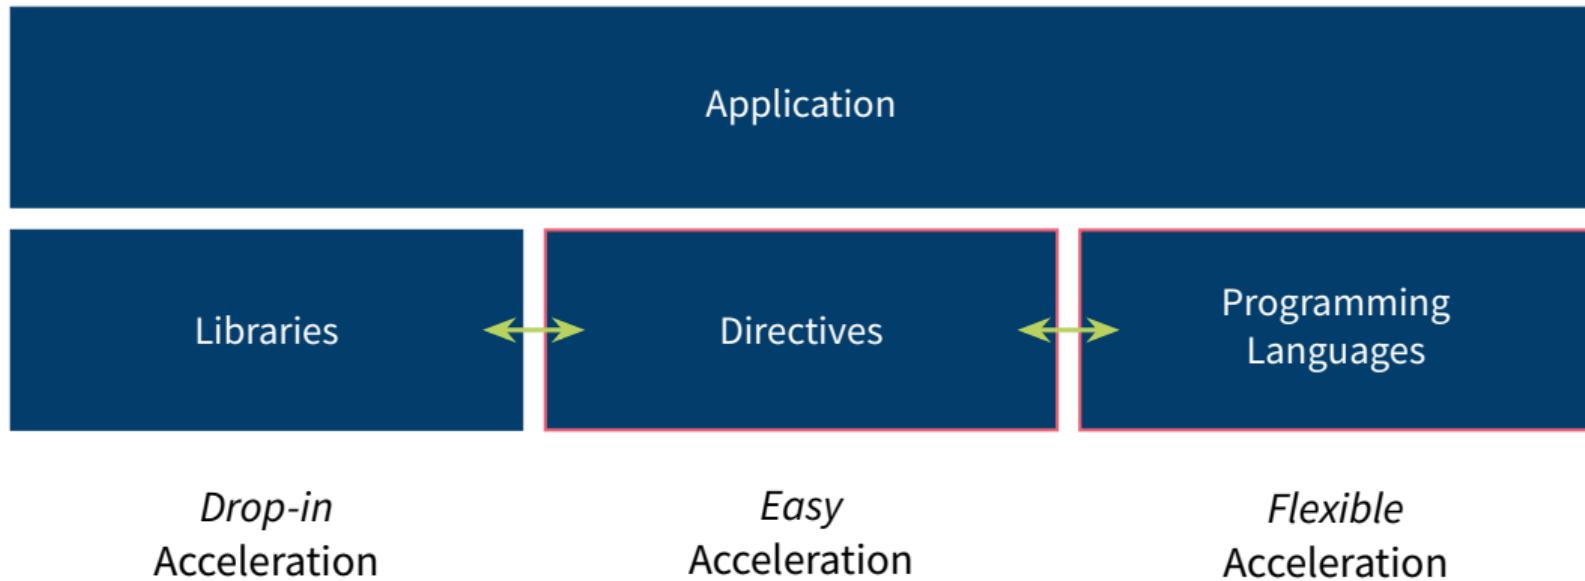


JÜLICH
SUPERCOMPUTING
CENTRE

Summary of Acceleration Possibilities



Summary of Acceleration Possibilities



⚠ Parallelism

Libraries are not enough?

You think you want to write your own GPU code?

Primer on Parallel Scaling

Amdahl's Law

Possible maximum speedup for
 N parallel processors

Total Time $t = t_{\text{serial}} + t_{\text{parallel}}$

Primer on Parallel Scaling

Amdahl's Law

Possible maximum speedup for
 N parallel processors

Total Time $t = t_{\text{serial}} + t_{\text{parallel}}$

N Processors $t(N) = t_s + t_p/N$

Primer on Parallel Scaling

Amdahl's Law

Possible maximum speedup for
 N parallel processors

Total Time $t = t_{\text{serial}} + t_{\text{parallel}}$

N Processors $t(N) = t_s + t_p/N$

Speedup $s(N) = t/t(N) = \frac{t_s+t_p}{t_s+t_p/N}$

Primer on Parallel Scaling

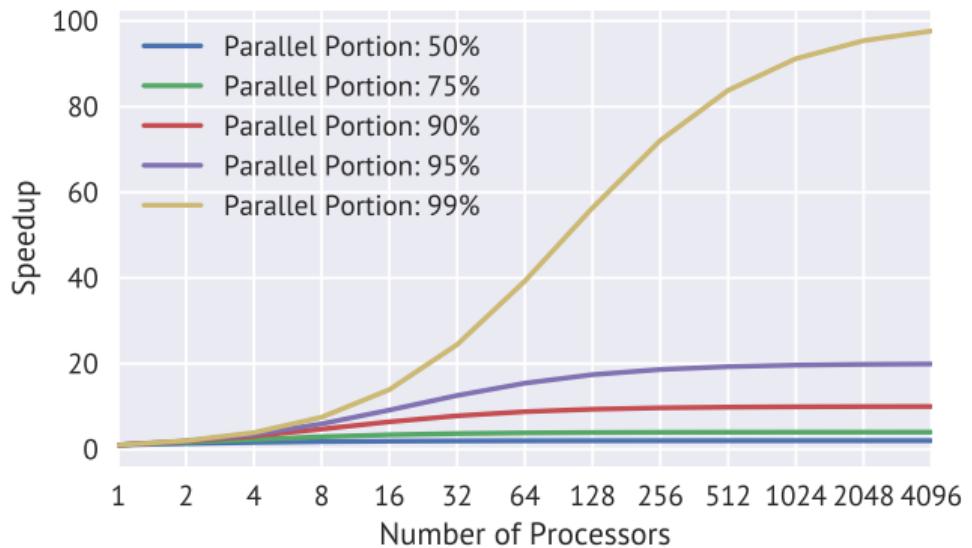
Amdahl's Law

Possible maximum speedup for N parallel processors

Total Time $t = t_{\text{serial}} + t_{\text{parallel}}$

N Processors $t(N) = t_s + t_p/N$

Speedup $s(N) = t/t(N) = \frac{t_s+t_p}{t_s+t_p/N}$



⚠ Parallelism

Parallel programming is not easy!

Things to consider:

- Is my application **computationally intensive** enough?
- What are the levels of **parallelism**?
- How much **data** needs to be **transferred**?
- Is the **gain** worth the **pain**?

Alternatives

The twilight

There are alternatives to CUDA C, which **can** ease the *pain*...

- OpenACC, OpenMP
- Thrust
- Kokkos, RAJA, ALPAKA, SYCL, DPC++, pSTL
- PyCUDA, Cupy, Numba

Other alternatives

- CUDA Fortran
- HIP
- OpenCL

Programming GPUs

CUDA C/C++

Preface: CPU

A simple CPU program!

SAXPY: $\vec{y} = a\vec{x} + \vec{y}$, with single precision

Part of LAPACK BLAS Level 1

```
void saxpy(int n, float a, float * x, float * y) {
    for (int i = 0; i < n; i++)
        y[i] = a * x[i] + y[i];
}

int a = 42;
int n = 10;
float x[n], y[n];
// fill x, y

saxpy(n, a, x, y);
```

CUDA SAXPY

With runtime-managed data transfers

```
__global__ void saxpy_cuda(int n, float a, float * x, float * y) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < n)
        y[i] = a * x[i] + y[i];
}

int a = 42;
int n = 10;
float x[n], y[n];
// fill x, y
cudaMallocManaged(&x, n * sizeof(float));
cudaMallocManaged(&y, n * sizeof(float));

saxpy_cuda<<<2, 5>>>(n, a, x, y);

cudaDeviceSynchronize();
```

CUDA SAXPY

With runtime-managed data transfers

```
__global__ void saxpy_cuda(int n, float a, float * x, float * y) {  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
    if (i < n)  
        y[i] = a * x[i] + y[i];  
}  
  
int a = 42;  
int n = 10;  
float x[n], y[n];  
// fill x, y  
cudaMallocManaged(&x, n * sizeof(float));  
cudaMallocManaged(&y, n * sizeof(float));  
  
saxpy_cuda<<<2, 5>>>(n, a, x, y);  
  
cudaDeviceSynchronize();
```

Specify kernel

ID variables

Guard against too many threads

Allocate GPU-capable memory

Call kernel
2 blocks, each 5 threads

Wait for kernel to finish

CUDA's Parallel Model

In software: Threads, Blocks

- Methods to exploit parallelism:

CUDA's Parallel Model

In software: Threads, Blocks

- Methods to exploit parallelism:

- Thread



CUDA's Parallel Model

In software: Threads, Blocks

- Methods to exploit parallelism:
 - Threads



CUDA's Parallel Model

In software: Threads, Blocks

- Methods to exploit parallelism:

- Threads → Block



CUDA's Parallel Model

In software: Threads, Blocks

- Methods to exploit parallelism:

- Threads → Block
- Block

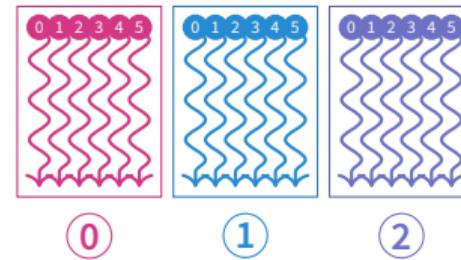


CUDA's Parallel Model

In software: Threads, Blocks

- Methods to exploit parallelism:

- Threads → Block
- Block

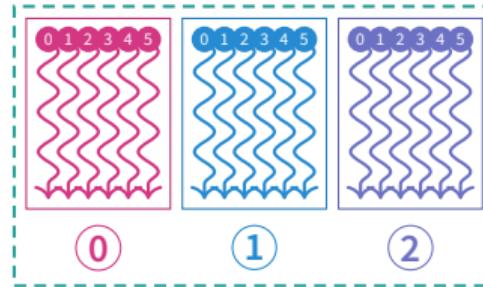


CUDA's Parallel Model

In software: Threads, Blocks

- Methods to exploit parallelism:

- Threads → Block
- Block → Grid

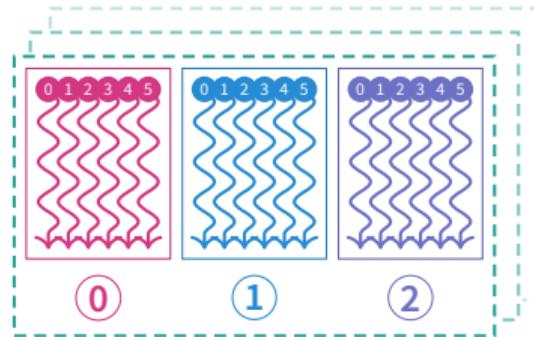


CUDA's Parallel Model

In software: Threads, Blocks

- Methods to exploit parallelism:

- Threads → Block
- Block → Grid
- Threads & blocks in 3D 



CUDA's Parallel Model

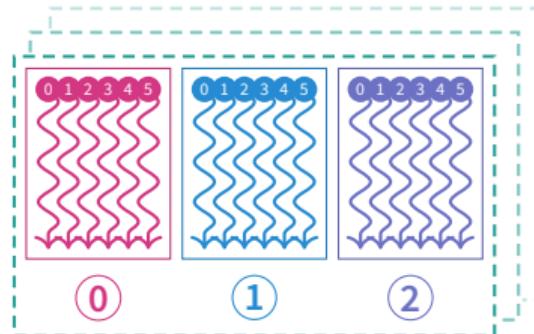
In software: Threads, Blocks

- Methods to exploit parallelism:

- Threads → Block

- Block → Grid

- Threads & blocks in 3D



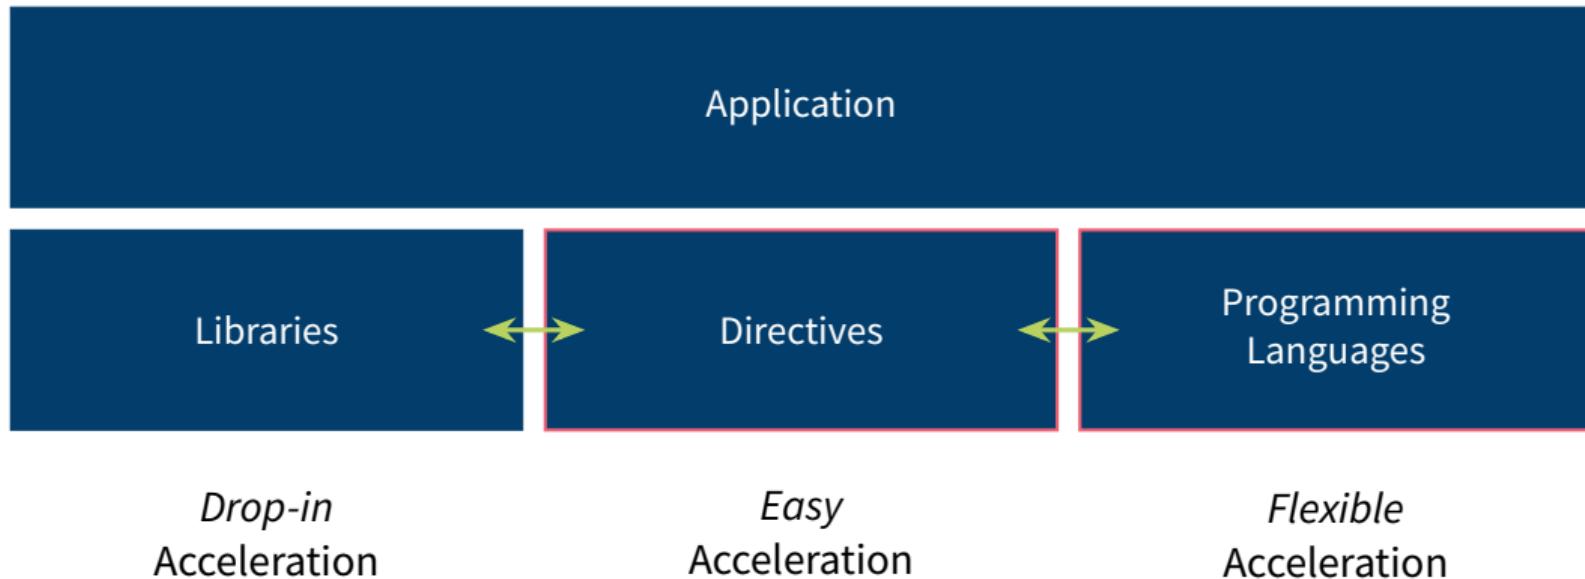
- Execution entity: **threads**

- Lightweight → fast switching!

- 1000s threads execute simultaneously → order non-deterministic!

- **OpenACC** takes care of threads and blocks for you!
→ Block configuration is just an optimization!

Summary of Acceleration Possibilities



Summary of Acceleration Possibilities



Conclusions

Conclusions

- GPUs achieve performance by specialized hardware → **threads**
 - Faster *time-to-solution*
 - Lower *energy-to-solution*
- GPU acceleration can be done by different means
- Libraries are the easiest, CUDA the fullest
- OpenACC good compromise

Conclusions

- GPUs achieve performance by specialized hardware → **threads**
 - Faster *time-to-solution*
 - Lower *energy-to-solution*
- GPU acceleration can be done by different means
- Libraries are the easiest, CUDA the fullest
- **OpenACC** good compromise

Conclusions

- GPUs achieve performance by specialized hardware → **threads**
 - Faster *time-to-solution*
 - Lower *energy-to-solution*
- GPU acceleration can be done by different means
- Libraries are the easiest, CUDA the fullest
- **OpenACC** good compromise

Thank you
for your attention!
a.herten@fz-juelich.de

Appendix

Appendix

Glossary

References

Glossary I

- AMD Manufacturer of **CPUs** and **GPUs**. [3](#), [4](#), [5](#), [6](#), [7](#), [8](#), [9](#)
- Ampere GPU architecture from **NVIDIA** (announced 2019). [13](#), [14](#), [15](#)
- API A programmatic interface to software by well-defined functions. Short for application programming interface. [107](#)
- ATI Canada-based **GPUs** manufacturing company; bought by AMD in 2006. [3](#), [4](#), [5](#), [6](#), [7](#), [8](#), [9](#)
- CUDA Computing platform for **GPUs** from NVIDIA. Provides, among others, CUDA C/C++. [2](#), [3](#), [4](#), [5](#), [6](#), [7](#), [8](#), [9](#), [84](#), [85](#), [87](#), [88](#), [89](#), [90](#), [91](#), [92](#), [93](#), [94](#), [95](#), [96](#), [97](#), [101](#), [102](#), [103](#), [107](#)
- JSC Jülich Supercomputing Centre, the supercomputing institute of Forschungszentrum Jülich, Germany. [107](#)

Glossary II

JURECA A multi-purpose supercomputer at JSC. [15](#)

JUWELS Jülich's new supercomputer, the successor of JUQUEEN. [12](#), [13](#), [14](#)

NVIDIA US technology company creating GPUs. [3](#), [4](#), [5](#), [6](#), [7](#), [8](#), [9](#), [12](#), [13](#), [14](#), [15](#), [58](#), [59](#), [60](#), [106](#), [108](#)

OpenACC Directive-based programming, primarily for many-core machines. [1](#), [84](#), [89](#), [90](#), [91](#), [92](#), [93](#), [94](#), [95](#), [96](#), [97](#)

OpenCL The *Open Computing Language*. Framework for writing code for heterogeneous architectures ([CPU](#), [GPU](#), DSP, FPGA). The alternative to [CUDA](#). [3](#), [4](#), [5](#), [6](#), [7](#), [8](#), [9](#), [84](#)

OpenGL The *Open Graphics Library*, an [API](#) for rendering graphics across different hardware architectures. [3](#), [4](#), [5](#), [6](#), [7](#), [8](#), [9](#)

OpenMP Directive-based programming, primarily for multi-threaded machines. [84](#)

Glossary III

SAXPY Single-precision $A \times X + Y$. A simple code example of scaling a vector and adding an offset. [86](#), [87](#), [88](#)

Tesla The [GPU](#) product line for general purpose computing computing of [NVIDIA](#). [12](#)

Thrust A parallel algorithms library for (among others) GPUs. See <https://thrust.github.io/>. [84](#)

CPU Central Processing Unit. [12](#), [15](#), [20](#), [21](#), [22](#), [26](#), [27](#), [28](#), [29](#), [30](#), [31](#), [32](#), [33](#), [34](#), [35](#), [36](#), [37](#), [51](#), [52](#), [53](#), [54](#), [55](#), [56](#), [57](#), [58](#), [59](#), [60](#), [86](#), [106](#), [107](#)

GPU Graphics Processing Unit. [1](#), [2](#), [3](#), [4](#), [5](#), [6](#), [7](#), [8](#), [9](#), [12](#), [13](#), [14](#), [15](#), [16](#), [17](#), [18](#), [20](#), [21](#), [22](#), [23](#), [24](#), [25](#), [26](#), [27](#), [28](#), [29](#), [30](#), [31](#), [32](#), [33](#), [34](#), [35](#), [36](#), [37](#), [38](#), [39](#), [41](#), [42](#), [51](#), [52](#), [53](#), [54](#), [55](#), [56](#), [57](#), [58](#), [59](#), [60](#), [64](#), [65](#), [66](#), [68](#), [71](#), [72](#), [73](#), [74](#), [75](#), [78](#), [85](#), [88](#), [101](#), [102](#), [103](#), [106](#), [107](#), [108](#)

Glossary IV

SIMD Single Instruction, Multiple Data. [51](#), [52](#), [53](#), [54](#), [55](#), [56](#), [57](#), [58](#), [59](#), [60](#)

SIMT Single Instruction, Multiple Threads. [23](#), [24](#), [25](#), [38](#), [39](#), [41](#), [42](#), [51](#), [52](#), [53](#), [54](#), [55](#),
[56](#), [57](#), [58](#), [59](#), [60](#)

SM Streaming Multiprocessor. [51](#), [52](#), [53](#), [54](#), [55](#), [56](#), [57](#), [58](#), [59](#), [60](#)

SMT Simultaneous Multithreading. [51](#), [52](#), [53](#), [54](#), [55](#), [56](#), [57](#), [58](#), [59](#), [60](#)

References I

- [2] Kenneth E. Hoff III et al. “Fast Computation of Generalized Voronoi Diagrams Using Graphics Hardware.” In: *Proceedings of the 26th Annual Conference on Computer Graphics and Interactive Techniques*. SIGGRAPH ’99. New York, NY, USA: ACM Press/Addison-Wesley Publishing Co., 1999, pp. 277–286. ISBN: 0-201-48560-5. DOI: [10.1145/311535.311567](https://doi.org/10.1145/311535.311567). URL: <http://dx.doi.org/10.1145/311535.311567> (pages 3–9).
- [3] Chris McClanahan. “History and Evolution of GPU Architecture.” In: *A Survey Paper* (2010). URL: <http://mcclanahoochie.com/blog/wp-content/uploads/2011/03/gpu-hist-paper.pdf> (pages 3–9).
- [4] Jack Dongarra et al. *TOP500*. June 2019. URL: <https://www.top500.org/lists/2019/06/> (pages 3–9).

References II

- [5] Jack Dongarra et al. *Green500*. June 2019. URL:
<https://www.top500.org/green500/lists/2019/06/> (pages 3–9).
- [6] Karl Rupp. *Pictures: CPU/GPU Performance Comparison*. URL:
<https://www.karlrupp.net/2013/06/cpu-gpu-and-mic-hardware-characteristics-over-time/> (pages 10, 11).
- [10] Wes Breazell. *Picture: Wizard*. URL:
<https://thenounproject.com/wes13/collection/its-a-wizards-world/> (pages 71–75).

References: Images, Graphics I

- [1] Héctor J. Rivas. *Color Reels*. Freely available at Unsplash. URL:
<https://unsplash.com/photos/87hFrPk3V-s>.
- [7] Forschungszentrum Jülich GmbH (Ralf-Uwe Limbach). *JUWELS Booster*.
- [8] Mark Lee. *Picture: kawasaki ninja*. URL:
<https://www.flickr.com/photos/pochacco20/39030210/> (pages 20, 21).
- [9] Shearings Holidays. *Picture: Shearings coach 636*. URL:
<https://www.flickr.com/photos/shearings/13583388025/> (pages 20, 21).